

ML Interview Take Home Assignment (HK)

👤 Created by	Aaron Chan
🕒 Created time	@December 18, 2025 11:04 PM
🏷️ Tags	

ML Interview Take-Home Assignment (HK)

Role: Generative AI Engineer

Company: finiti.legal

1. Overview

As an **AI Engineer**, you will build a **prototype Retrieval-Augmented Generation (RAG) assistant** for a legal use case: **assisting securities lawyers and finance professionals in drafting and reviewing SEC 10-K disclosures**.

This assignment is intentionally **end-to-end**:

- Document ingestion
- Retrieval + generation
- Interactive clarification of missing financial inputs
- Clean, usable API or CLI interface

The goal is not UI polish, but **correctness, usability, and domain awareness**.

2. Input Data

2.1 Primary Documents

You will work with past one year's **HTML 10-K filings from the SEC EDGAR website** for the **and business from the user for** following companies:

- NVDA
- MSFT
- KO
- NKE
- AMZN
- DASH
- TJX
- DRI

Example 10-K filing (UBER):

[https://www.sec.gov/ix?
doc=/Archives/edgar/data/0001543151/000154315125000008/uber-20241231.htm](https://www.sec.gov/ix?doc=/Archives/edgar/data/0001543151/000154315125000008/uber-20241231.htm)

Helpful links:

- SEC EDGAR search: <https://www.sec.gov/search-filings>
 - Public investor-relations pages for company background
-

2.2 Time Scope

- Use **information relevant to the most recent fiscal year**
 - Example: To generate a **2024 10-K narrative**, rely on **2023 financial and business data**
 - Do **not** hallucinate future or unpublished information
-

3. Output Requirements

You must design an **interactive assistant** that produces draft narratives for:

1. **Item 1. Business**
2. **Item 7. Management's Discussion & Analysis (MD&A)**

Your system must:

- Use **RAG** over the provided filings

- Generate **clear, legally-appropriate narrative text**
 - Be usable by **non-technical securities lawyers or finance professionals**
-

4. Critical Requirement: Interactive Financial Data Collection

Why This Matters

In real legal workflows, the structure or the shell of the **MD&A and other sections are created based on prior or competitor company filings**. Then business and financial data are gathered from the users and “filled in” to generate the current filing.

It often requires **current-year financial inputs** that may not yet exist in structured documents.

What Your System Must Do

Your assistant must be able to:

1. Identify missing or required financial inputs needed to complete MD&A

Examples include:

- Revenue and revenue growth
- Year-over-year or quarter-over-quarter comparisons
- Segment-level performance
- Operating income, margins, cash flow
- Material one-time events or cost changes

2. Ask the user targeted follow-up questions when inputs are missing or when the model is not confident with the extracted data

- Questions must be clear and business-friendly
- Avoid technical or ML-specific language

3. Accept user-provided financial data via chat (UI or CLI) in **any format**:

- Standalone numbers
- Markdown tables
- HTML tables

4. Incorporate the provided data into the MD&A narrative

- Clearly grounded in user input
- Written in appropriate 10-K tone and structure

You may assume the data provided by the user is accurate. Validation is optional.

5. Evaluation Examples (What “Good” Looks Like)

5.1 Sample Prompt — UBER 10-K Generation

User Prompt

Generate the Business and MD&A sections for UBER's 2024 Form 10-K.

5.2 Expected System Behavior

Step 1: Retrieval & Initial Draft

- Retrieve relevant sections from UBER's prior year's (e.g., 2023) 10-K filing
- Generate a **Business** section draft for the current year (e.g., 2024) using RAG

Step 2: Identify Missing Inputs for MD&A

The system should recognize missing financial data and respond with something like

To update the MD&A section, I need the following information for fiscal year 2024:

- Total revenue and year-over-year growth
- Key drivers of revenue changes
- Operating income or loss

- Notable cost or margin changes
- Any material one-time events

You can provide this data as numbers, a table, or pasted financial statements.

5.3 Sample User Financial Input

Metric	FY 2024	FY 2023
Revenue	\$37.2B	\$33.9B
Revenue Growth	9.7%	16.9%
Operating Income	\$1.9B	\$1.1B
Free Cash Flow	\$3.4B	\$3.1B

5.4 Expected Output Formatting

Your final output should be **clearly structured**, for example:

Item 1. Business

Concise narrative summary grounded in retrieved 10-K content.

No hallucinations. No unsupported claims.

Item 7. Management's Discussion and Analysis of Financial Condition and Results of Operations

Narrative that:

- Explicitly reflects user-provided financial data
- Explains drivers of performance
- Matches the tone and structure of a real 10-K MD&A

6. Tech Stack

- Python

- LLM frameworks (any of the following or equivalents):
 - LangChain
 - LangGraph
 - AutoGen
 - PydanticAI
 - CrewAI
- Vector database of your choice (e.g., Pinecone, Weaviate, Qdrant, FAISS)
- **FastAPI** for backend endpoints

No frontend is required.

CLI-based demos are acceptable.

| You are not evaluated on your choice of LLM or framework.

7. User Persona

Assume your users are:

- Securities lawyers
- Finance professionals

They are comfortable with filings but **not** with AI or engineering details.

Your system should:

- Ask clear, structured questions
 - Avoid technical jargon
 - **[Important] Feel like a human legal assistant, not hard coded chat**
-

8. Evaluation Criteria

Category	What We Look For
RAG Design	Correct retrieval, clean chunking, grounding

Category	What We Look For
Interactive Logic	Correctly identifying missing MD&A inputs
Output Quality	Accuracy, clarity, legal tone
Usability	Clear prompts, sensible defaults
Performance	Reasonable latency and system design

9. Optional Enhancements (Not Required)

- Source citations per paragraph
 - Confidence indicators for generated text
 - Simple year-over-year comparison logic
 - Audit-friendly logging of user inputs vs output
-

Final Note

This assignment reflects how finiti.legal is used in practice:

AI-assisted drafting with human-provided financial, business and operational inputs, strong grounding, and lawyer-friendly UX.

Examples:

Financial inputs

- 2024 revenue by segment (Mobility, Delivery, Freight) and key YoY % changes.
- Adjusted EBITDA, net income/loss, cash and debt balances, major capex amounts.

Business Inputs

- New products or services launched and any lines that were discontinued.
- Entry into or exit from key countries, cities, or verticals; major partnerships or JVs.

Operational Inputs

- Changes in driver/merchant policies, pricing model, or incentive structures.
- Key operational risks or events (strikes, outages, regulatory actions, safety incidents).

Over-engineering is not rewarded.

Clarity, correctness, and domain fit are.